

Exp:7**Implement Linear and Logistic Regression**

a)Linear regression

Sample data

```
heights <- c(150, 160, 165, 170, 175, 180, 185)
```

```
weights <- c(55, 60, 62, 68, 70, 75, 80)
```

Create a data frame

```
data <- data.frame(heights, weights)
```

Fit a linear regression model

```
linear_model <- lm(weights ~ heights, data = data)
```

Print the summary of the model

```
print(summary(linear_model))
```

Plotting the data and regression line

```
plot(data$heights, data$weights,
```

```
main = "Linear Regression: Weight vs. Height",
```

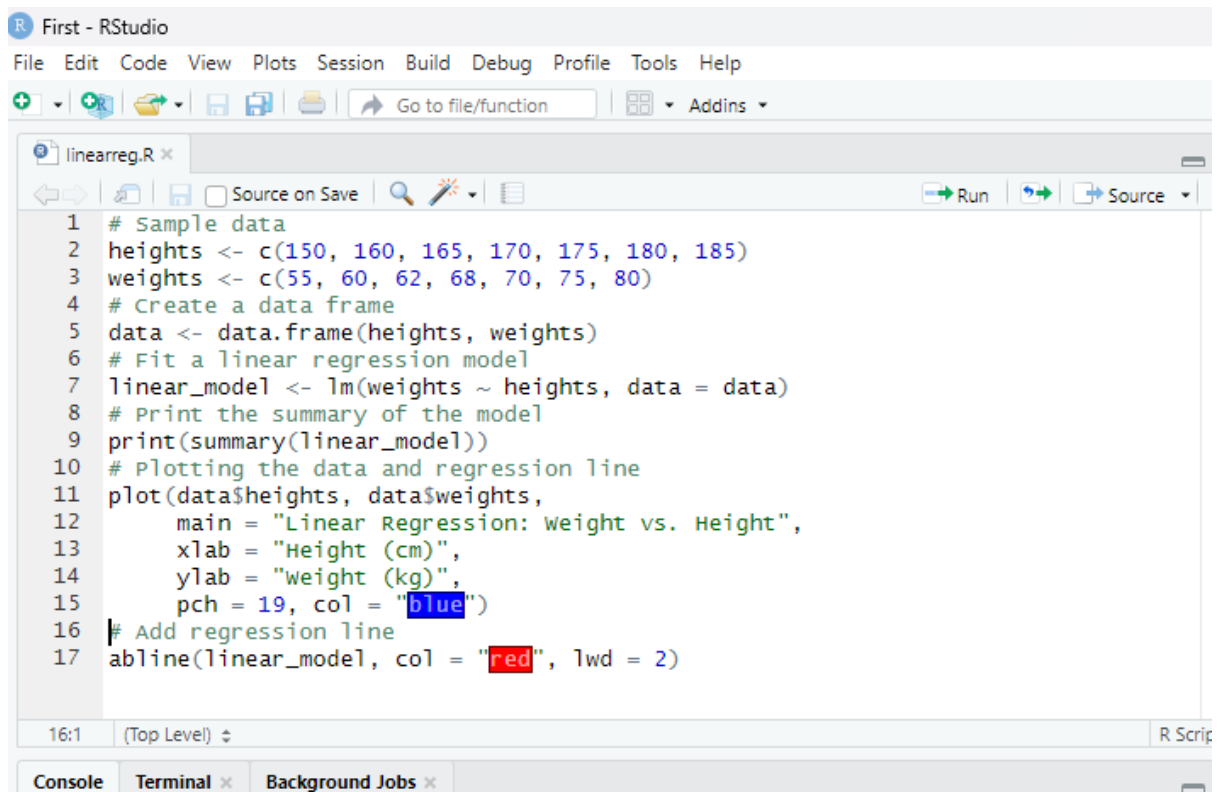
```
xlab = "Height (cm)",
```


```
ylab = "Weight (kg)",
```

```
pch = 19, col = "blue")
```

Add regression line

```
abline(linear_model, col = "red", lwd = 2)
```





```
# Load the dataset
```

```
# Convert 'am' to a factor (categorical variable)
```

```
# Fit a logistic regression model
```

```
# Print the summary of the model
```

```
print(summary(logistic_model))

# Predict probabilities for the logistic model
predicted_probs <- predict(logistic_model, type = "response")

# Display the predicted probabilities
print(predicted_probs)

# Plotting the data and logistic regression curve
plot(mtcars$mpg, as.numeric(mtcars$am) - 1,
     main = "Logistic Regression: Transmission vs. MPG",
     xlab = "Miles Per Gallon (mpg)",
     ylab = "Probability of Manual Transmission",
     pch = 19, col = "blue")

# Add the logistic regression curve
curve(predict(logistic_model, data.frame(mpg = x), type = "response"),
      add = TRUE, col = "red", lwd = 2)
```

The screenshot displays the RStudio environment with a script editor and a console window.

Script Editor (linearreg.R):

```

1 # Load the dataset
2 data(mtcars)
3 # Convert 'am' to a factor (categorical variable)
4 mtcars$am <- factor(mtcars$am, levels = c(0, 1), labels = c("Automatic", "Manual"))
5 # Fit a logistic regression model
6 logistic_model <- glm(am ~ mpg, data = mtcars, family = binomial)
7 # Print the summary of the model
8 print(summary(logistic_model))
9 # Predict probabilities for the logistic model
10 predicted_probs <- predict(logistic_model, type = "response")
11 # Display the predicted probabilities
12 print(predicted_probs)
13 # Plotting the data and logistic regression curve
14 plot(mtcars$mpg, as.numeric(mtcars$am) - 1,
15      main = "Logistic Regression: Transmission vs. MPG",
16      xlab = "Miles Per Gallon (mpg)",
17      ylab = "Probability of Manual Transmission",
18      nch = 19, col = "blue")
19

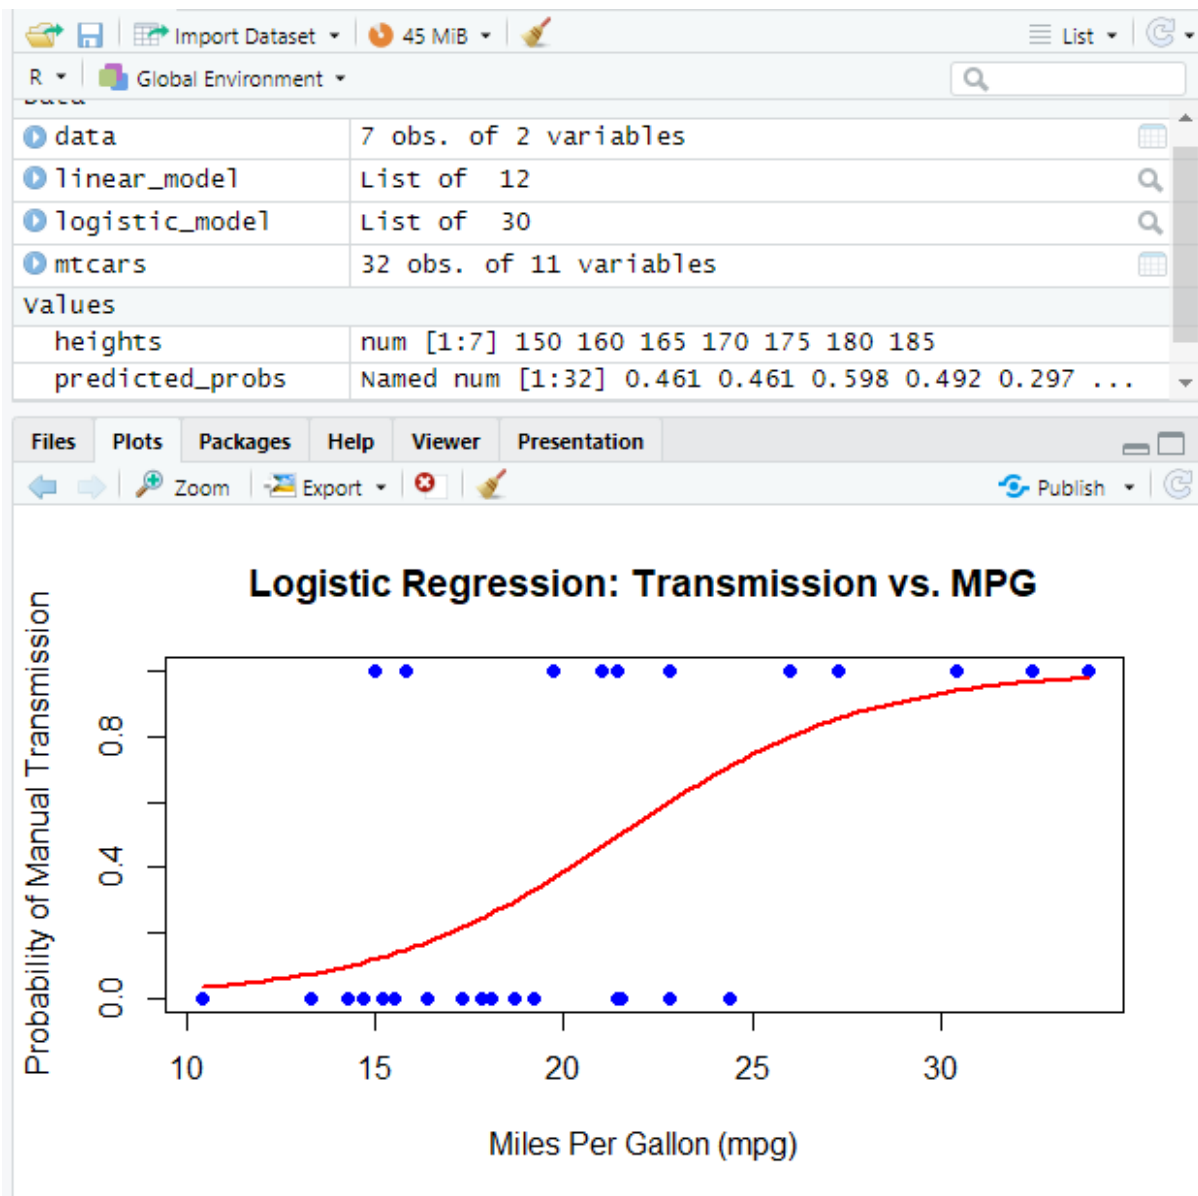
```

Console Output:

```

R 4.4.1 - C:/R/First/
0.17246396
0.59789839
0.21552479
Chrysler Imperial
0.11005178
Toyota Corona
0.49939484
Pontiac Firebird
0.32991148
Ford Pantera L
0.14773451
0.32991148
0.12601104
0.96591395
Dodge Challenger
0.13650937
Fiat X1-9
0.85549212
Ferrari Dino
0.36468861
0.24260966
0.03197098
Honda Civic
0.93878132
AMC Javelin
0.12601104
Porsche 914-2
0.79886349
Maserati Bora
0.11940215
0.17246396
0.17246396
Lincoln Continental
0.03197098
Toyota Corolla
0.97821971
Camaro Z28
0.07446438
Lotus Europa
0.93878132
Volvo 142E
0.49171990

```



Exp8: Implement SVM/Decision tree classification techniques

a) SVM IN R

```
# Install and load the e1071 package (if not already installed)

install.packages("e1071")

library(e1071)

# Load the iris dataset

data(iris)

# Inspect the first few rows of the dataset

head(iris)

# Split the data into training (70%) and testing (30%) sets

set.seed(123) # For reproducibility

sample_indices <- sample(1:nrow(iris), 0.7 * nrow(iris))

train_data <- iris[sample_indices, ]

test_data <- iris[-sample_indices, ]

# Fit the SVM model

svm_model <- svm(Species ~ ., data = train_data, kernel = "radial")

# Print the summary of the model

summary(svm_model)

# Predict the test set

predictions <- predict(svm_model, newdata = test_data)

# Evaluate the model's performance

confusion_matrix <- table(Predicted = predictions, Actual = test_data$Species)

print(confusion_matrix)

# Calculate accuracy

accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)

cat("Accuracy:", accuracy * 100, "%\n")
```

The screenshot shows an R Studio window with a script editor and a console. The script editor contains the following R code:

```
1 # Install and load the e1071 package (if not already installed)
2 install.packages("e1071")
3 library(e1071)
4 # Load the iris dataset
5 data(iris)
6 # Inspect the first few rows of the dataset
7 head(iris)
8 # Split the data into training (70%) and testing (30%) sets
9 set.seed(123) # For reproducibility
10 sample_indices <- sample(1:nrow(iris), 0.7 * nrow(iris))
11 train_data <- iris[sample_indices, ]
12 test_data <- iris[-sample_indices, ]
13 # Fit the SVM model
14 svm_model <- svm(Species ~ ., data = train_data, kernel = "radial")
15 # Print the summary of the model
16 summary(svm_model)
17 # Predict the test set
18 predictions <- predict(svm_model, newdata = test_data)
19 # Evaluate the model's performance
```

The console output shows the following messages:

```
R 4.4.1 - C:/R/First/
package 'proxy' successfully unpacked and MD5 sums checked
package 'e1071' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
  C:\Users\hp\AppData\Local\Temp\Rtmpe05wqE\downloaded_packages
Actual
Predicted   setosa versicolor virginica
setosa      14         0         0
versicolor  0         17         0
virginica   0          1        13
Accuracy: 97.77778 %
> |
```

b) Decision tree in R

Install and load the rpart package (if not already installed)

```
install.packages("rpart")
```

```
library(rpart)
```

Load the iris dataset

```
data(iris)
```

Split the data into training (70%) and testing (30%) sets

```
set.seed(123) # For reproducibility
```

```
sample_indices <- sample(1:nrow(iris), 0.7 * nrow(iris))
```

```
train_data <- iris[sample_indices, ]
```



```
test_data <- iris[-sample_indices, ]

# Fit the Decision Tree model

tree_model <- rpart(Species ~ ., data = train_data, method = "class")

# Print the summary of the model

summary(tree_model)

# Plot the Decision Tree

plot(tree_model)

text(tree_model, pretty = 0)

# Predict the test set

predictions <- predict(tree_model, newdata = test_data, type = "class")

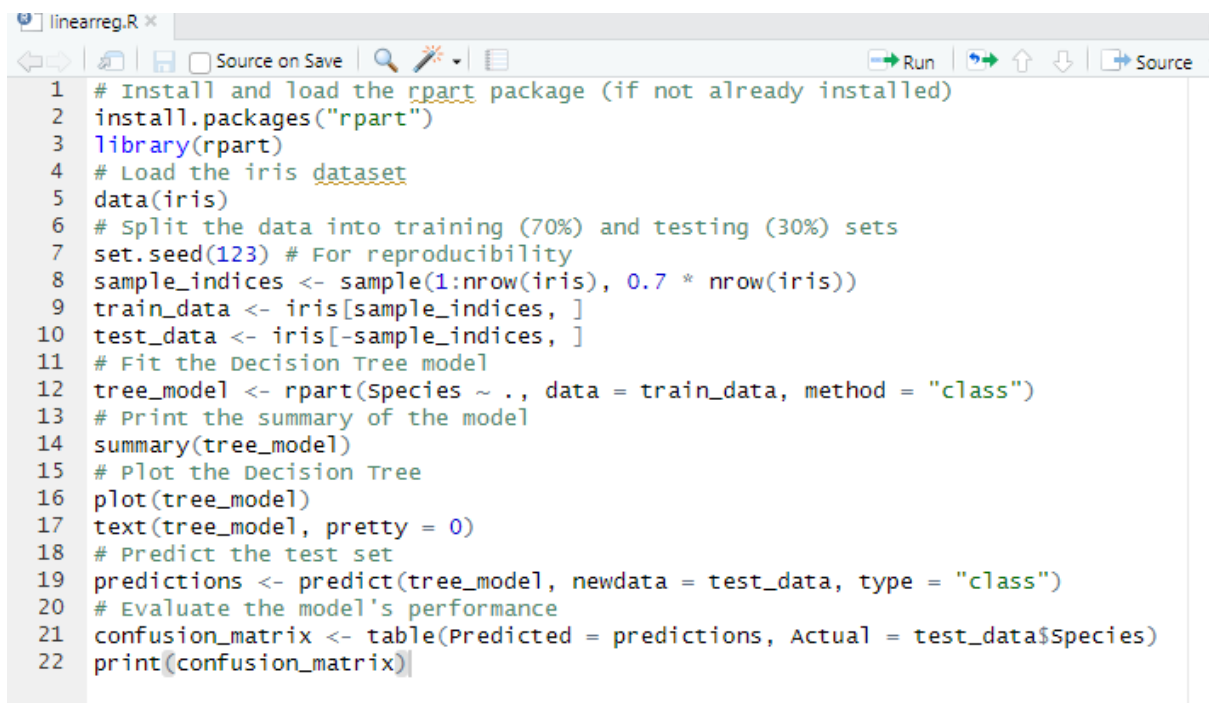
# Evaluate the model's performance

confusion_matrix <- table(Predicted = predictions, Actual = test_data$Species)

print(confusion_matrix)

accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)

cat("Accuracy:", accuracy * 100, "%\n")
```



```
1 # Install and load the rpart package (if not already installed)
2 install.packages("rpart")
3 library(rpart)
4 # Load the iris dataset
5 data(iris)
6 # Split the data into training (70%) and testing (30%) sets
7 set.seed(123) # For reproducibility
8 sample_indices <- sample(1:nrow(iris), 0.7 * nrow(iris))
9 train_data <- iris[sample_indices, ]
10 test_data <- iris[-sample_indices, ]
11 # Fit the Decision Tree model
12 tree_model <- rpart(Species ~ ., data = train_data, method = "class")
13 # Print the summary of the model
14 summary(tree_model)
15 # Plot the Decision Tree
16 plot(tree_model)
17 text(tree_model, pretty = 0)
18 # Predict the test set
19 predictions <- predict(tree_model, newdata = test_data, type = "class")
20 # Evaluate the model's performance
21 confusion_matrix <- table(Predicted = predictions, Actual = test_data$Species)
22 print(confusion_matrix)
```

package 'rpart' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\hp\AppData\Local\Temp\Rtmpe05wqe\downloaded_packages

Call:

```
rpart(formula = Species ~ ., data = train_data, method = "class")
n= 105
```

| | CP | nsplit | rel error | xerror | xstd |
|---|-----------|--------|------------|-----------|------------|
| 1 | 0.5294118 | 0 | 1.00000000 | 1.2058824 | 0.06232572 |
| 2 | 0.3970588 | 1 | 0.47058824 | 0.5441176 | 0.07198662 |
| 3 | 0.0100000 | 2 | 0.07352941 | 0.1176471 | 0.03997857 |

Variable importance

| Petal.width | Petal.Length | Sepal.Length | Sepal.width |
|-------------|--------------|--------------|-------------|
| 34 | 32 | 21 | 13 |

Node number 1: 105 observations, complexity param=0.5294118

predicted class=virginica expected loss=0.647619 P(node) =1

class counts: 36 32 37

probabilities: 0.343 0.305 0.352

left son=2 (36 obs) right son=3 (69 obs)

Primary splits:

Petal.Length < 2.45 to the left, improve=35.54783, (0 missing)

Petal.width < 0.8 to the left, improve=35.54783, (0 missing)

Sepal.Length < 5.45 to the left, improve=24.79179, (0 missing)

Sepal.width < 3.25 to the right, improve=12.34670, (0 missing)

Surrogate splits:

Petal.width < 0.8 to the left, agree=1.000, adj=1.000, (0 split)

Sepal.Length < 5.45 to the left, agree=0.924, adj=0.778, (0 split)

Sepal.width < 3.25 to the right, agree=0.819, adj=0.472, (0 split)

Node number 2: 36 observations

predicted class=setosa expected loss=0 P(node) =0.3428571

class counts: 36 0 0

probabilities: 1.000 0.000 0.000

Node number 3: 69 observations, complexity param=0.3970588

predicted class=virginica expected loss=0.4637681 P(node) =0.6571429

class counts: 0 32 37

probabilities: 0.000 0.464 0.536

left son=6 (35 obs) right son=7 (34 obs)

Primary splits:

Petal.width < 1.75 to the left, improve=25.291950, (0 missing)

Petal.Length < 4.75 to the left, improve=25.187810, (0 missing)

Sepal.Length < 6.15 to the left, improve= 5.974246, (0 missing)

Sepal.width < 2.45 to the left, improve= 2.411006, (0 missing)

Surrogate splits:

Petal.Length < 4.75 to the left, agree=0.913, adj=0.824, (0 split)

Sepal.Length < 6.15 to the left, agree=0.696, adj=0.382, (0 split)

Sepal.width < 2.65 to the left, agree=0.638, adj=0.265, (0 split)

Node number 6: 35 observations

predicted class=versicolor expected loss=0.1142857 P(node) =0.3333333

class counts: 0 31 4

probabilities: 0.000 0.886 0.114


Node number 7: 34 observations

predicted class=virginica expected loss=0.02941176 P(node) =0.3238095

class counts: 0 1 33

probabilities: 0.000 0.029 0.971

| | Actual | | |
|------------|--------|------------|-----------|
| Predicted | setosa | versicolor | virginica |
| setosa | 14 | 0 | 0 |
| versicolor | 0 | 18 | 1 |
| virginica | 0 | 0 | 12 |



Exp:9 Implement clustering techniques – Hierarchical and K-Means

a) HIERARCHIAL CLUSTERING

```
# Load the iris dataset
```

```
data(iris)
```

```
# Use only the numeric columns for clustering (exclude the Species column)
```

```
iris_data <- iris[, -5]
```

```
# Standardize the data
```

```
iris_scaled <- scale(iris_data)
```

```
# Compute the distance matrix
```

```
distance_matrix <- dist(iris_scaled, method = "euclidean")
```

```
# Perform hierarchical clustering using the "complete" linkage method
```

```
hc_complete <- hclust(distance_matrix, method = "complete")
```

```
# Plot the dendrogram
```

```
plot(hc_complete, main = "Hierarchical Clustering Dendrogram", xlab = "", sub = "", cex =  
0.6)
```

```
# Cut the tree to form 3 clusters
```

```
clusters <- cutree(hc_complete, k = 3)
```

```
# Print the cluster memberships
```

```
print(clusters)
```

```
# Add the clusters to the original dataset
```

```
iris$Cluster <- as.factor(clusters)
```

```
# Display the first few rows of the updated dataset
```

```
head(iris)
```



```
# Load the iris dataset

data(iris)

# Use only the numeric columns for clustering (exclude the Species column)

iris_data <- iris[, -5]


# Standardize the data

iris_scaled <- scale(iris_data)

# Set the number of clusters

set.seed(123) # For reproducibility

k <- 3 # Number of clusters

# Perform K-Means clustering

kmeans_result <- kmeans(iris_scaled, centers = k, nstart = 25)

# Print the K-Means result

print(kmeans_result)

# Print the cluster centers

print(kmeans_result$centers)

# Add the cluster assignments to the original dataset

iris$Cluster <- as.factor(kmeans_result$cluster)

# Display the first few rows of the updated dataset

head(iris)


# Plot the clusters

library(ggplot2)

ggplot(iris, aes(x = Sepal.Length, y = Sepal.Width, color = Cluster)) +
  geom_point(size = 3) +
  labs(title = "K-Means Clustering of Iris Dataset", x = "Sepal Length", y = "Sepal Width")
```

[illegible]

K-Means Clustering of Iris Dataset



Exp:10**VISUALIZE DATA USING ANY PLOTTING FRAMEWORK****1) SCATTER PLOT**

Install ggplot2 (if not already installed)

```
install.packages("ggplot2")
```

Load the ggplot2 package

```
library(ggplot2)
```

Scatter plot of Sepal.Length vs Sepal.Width, colored by Species

```
ggplot(data = iris, aes(x = Sepal.Length, y = Sepal.Width, color = Species)) +
```

```
geom_point(size = 3) + # Adds points
```

```
labs(title = "Scatter Plot of Sepal Dimensions",
```

```
x = "Sepal Length (cm)",
```

```
y = "Sepal Width (cm)") + # Adds axis labels and title
```

```
theme_minimal() # Applies a minimal theme
```

```
https://cran.rstudio.com/bin/windows/Rtools/
```

```
warning in install.packages :
```

```
package 'ggplot2' is in use and will not be installed
```

```
> source("C:/R/First/linearreg.R")
```

```
> # scatter plot of Sepal.Length vs Sepal.Width, colored by species
```

```
> ggplot(data = iris, aes(x = Sepal.Length, y = Sepal.Width, color = Species)) +
```

```
+ geom_point(size = 3) + # Adds points
```


```
+ labs(title = "Scatter Plot of Sepal Dimensions",
```

```
+ x = "Sepal Length (cm)",
```

```
+ y = "Sepal width (cm)") + # Adds axis labels and title
```

```
+ theme_minimal() # Applies a minimal theme
```

```
> |
```



2) BAR CHART

Install ggplot2 (if not already installed)

```
install.packages("ggplot2")
```

Load the ggplot2 package

```
library(ggplot2)
```

Bar plot of Species counts

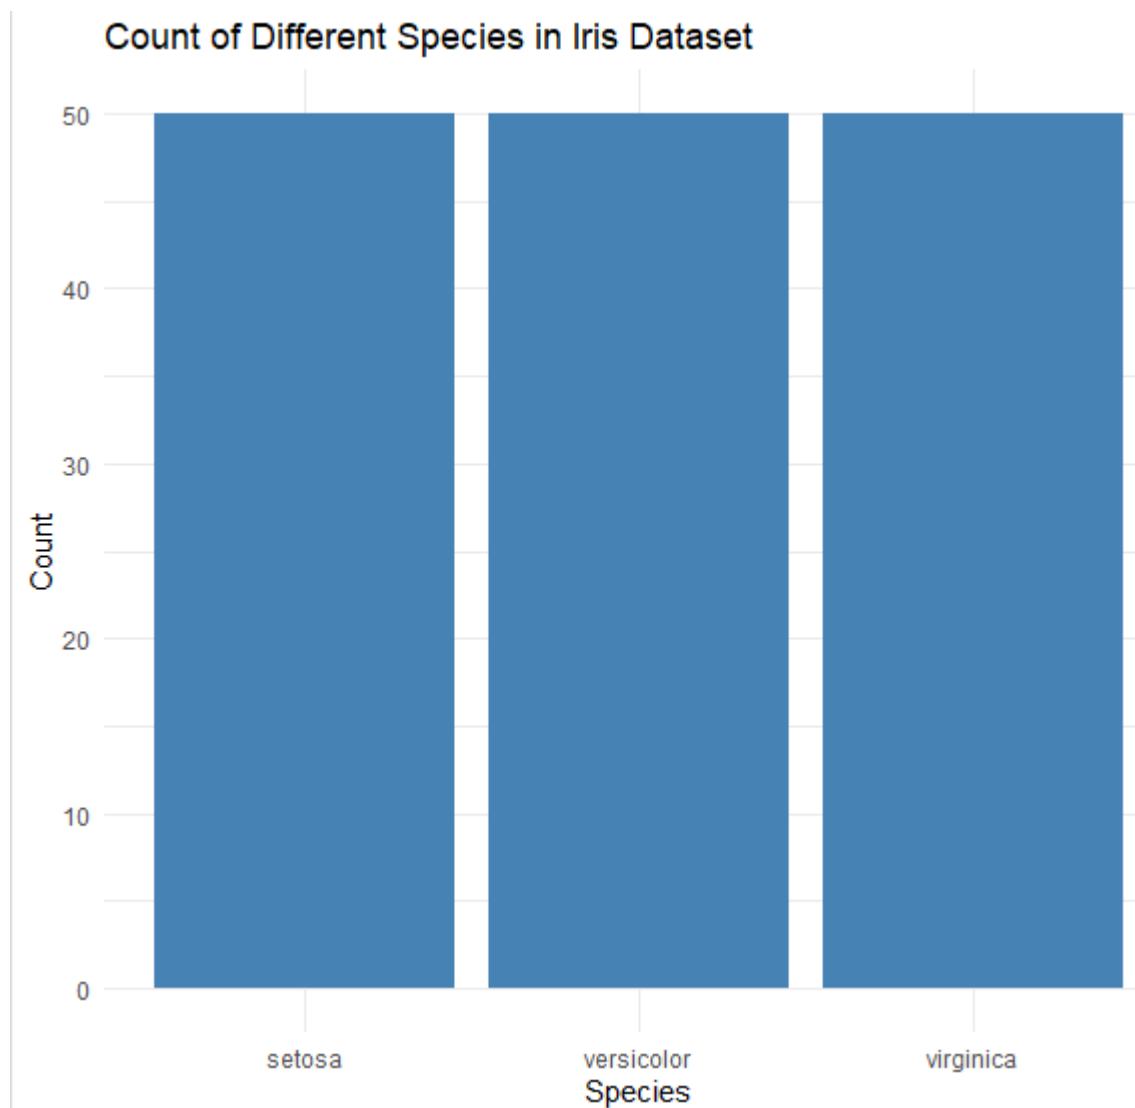
```
ggplot(data = iris, aes(x = Species)) +
```

```
geom_bar(fill = "steelblue") + # Adds bars filled with steel blue color
```

theme_minimal()

```
> source("C:/R/First/linearreg.R")
Error in install.packages : Updating loaded packages
> install.packages("ggplot2")
WARNING: Rtools is required to build R packages but is not currently installed. Please download and install the appropriate version of Rtools before proceeding:

https://cran.rstudio.com/bin/windows/Rtools/
Warning in install.packages :
  package 'ggplot2' is in use and will not be installed
> source("C:/R/First/linearreg.R")
> # Scatter plot of Sepal.Length vs Sepal.width, colored by species
> ggplot(data = iris, aes(x = Sepal.Length, y = Sepal.width, color = species)) +
+   geom_point(size = 3) + # Adds points
+   labs(title = "Scatter Plot of Sepal Dimensions",
+         x = "Sepal Length (cm)",
+         y = "Sepal width (cm)") + # Adds axis labels and title
+   theme_minimal() # Applies a minimal theme
> ggplot(data = iris, aes(x = Species)) +
+   geom_bar(fill = "steelblue") + # Adds bars filled with steel blue color
+   labs(title = "Count of Different Species in Iris Dataset",
+         x = "Species",
+         y = "Count") +
+   theme_minimal()
```



3) HISTOGRAM

Install ggplot2 (if not already installed)

```
install.packages("ggplot2")
```

Load the ggplot2 package

```
library(ggplot2)
```


Histogram of Sepal Length

```
ggplot(data = iris, aes(x = Sepal.Length)) +
```

```
geom_histogram(binwidth = 0.3, fill = "orange", color = "black") + # Adds
```

```
histogram bars
```

```
labs(title = "Histogram of Sepal Length",  
x = "Sepal Length (cm)",  
y = "Frequency") +  
theme_minimal()  
# Histogram of Sepal Length  
ggplot(data = iris, aes(x = Sepal.Length)) +  
  geom_histogram(binwidth = 0.3, fill = "orange", color = "black")  
labs(title = "Histogram of Sepal Length",  
x = "Sepal Length (cm)",  
y = "Frequency") +  
  theme_minimal()
```



4)BOX PLOT

Install ggplot2 (if not already installed)

```
install.packages("ggplot2")
```

Load the ggplot2 package

```
library(ggplot2)
```

Box plot of Sepal Length for each Species

```
ggplot(data = iris, aes(x = Species, y = Sepal.Length, fill = Species)) +
```

```
geom_boxplot() + # Adds box plot
```

```
labs(title = "Box Plot of Sepal Length by Species",
```

```
x = "Species",
```

```
y = "Sepal Length (cm)") +
```

```
theme_minimal()
```

```
ggplot(data = iris, aes(x = species, y = Sepal.Length, fill = species)) +  
  geom_boxplot() + # Adds box plot  
  labs(title = "Box Plot of Sepal Length by Species",  
        x = "species",  
        y = "Sepal Length (cm)") +  
  theme_minimal()
```

